

Setup

- Get Chapel on Peano, then try a program

○ ○ ○

```
<ssh into peano>  
$ module load git singularity gnu8  
$ git clone git@github.com:ghbrown/chapel-introduction.git  
$ source /tmp/chapel/get_chpl.sh
```



○ ○ ○

```
$ cd chapel-introduction/src  
$ chpl hello.chpl  
$ ./hello
```

- or

○ ○ ○

```
<on your local machine>  
$ browser https://github.com/ghbrown/chapel-introduction/tree/main  
$ browser https://shorturl.at/ezIV7
```

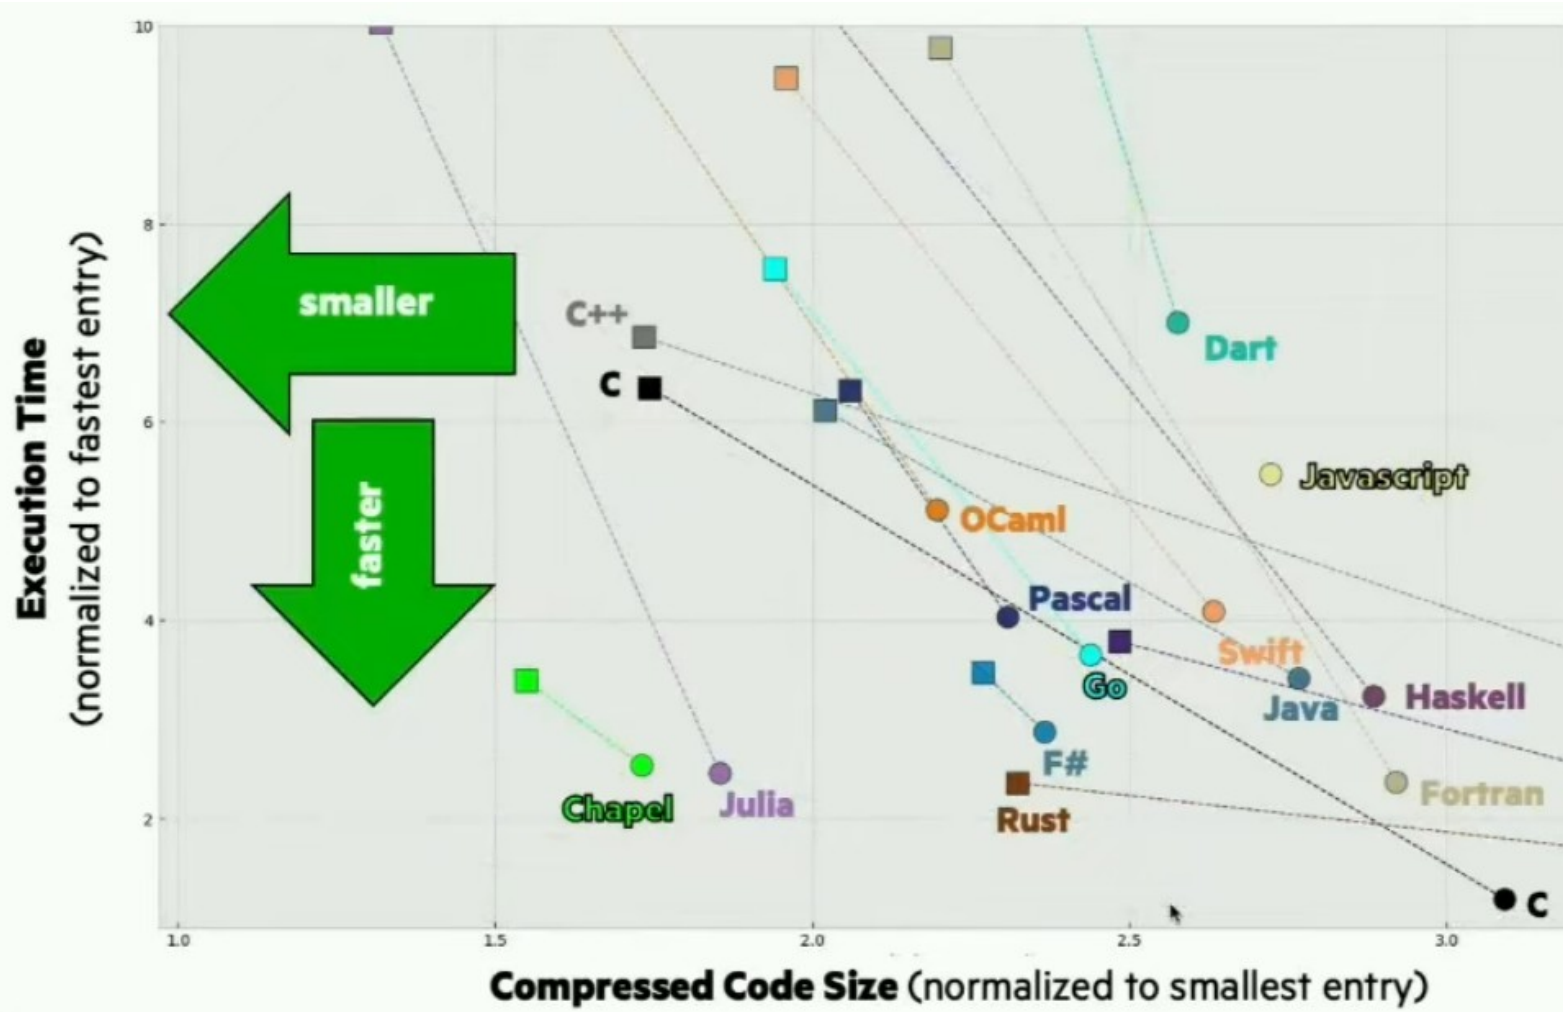


Chapel Parallel Programming Language: a **user** introduction

History

- Cray's entry into DARPA's High Productivity Computing Systems project
 - “there exists a critical need for improved software tools, standards, and methodologies for effective utilization of multiprocessor computers” – Lusk and Yelick, 2007
- A language for high performance and productivity parallel programming
- Development now continued at HPE

Motivation: language benchmarks



Chapel: serial language features

- static typing (with type inference)

○ ○ ○

```
var s : string = 'foo';
```

```
var a : real = 1.0;
```

```
var b : real = 3.0;
```

```
var c : real = a + b;
```

○ ○ ○

```
var s = 'foo';
```

```
var a = 1.0;
```

```
var b = 3.0;
```

```
var c = a + b;
```

Chapel: serial language features

- generics
 - functions can work for multiple input types
 - not present in C or Fortran (without interfaces)

○ ○ ○

```
proc square(v) {  
    return v*v;  
}
```

```
var a : real(32) = 1.0;
```

```
var b : real(64) = 2.0;
```

```
writeln(square(a));
```

```
writeln(square(b));
```

Chapel: serial language features

- first class arrays (with fancy indexing)

○○○

```
proc frobenius_norm(A) {  
    return sqrt(+ reduce A**2); // sqrt(sum of a_ij^2 for all i,j)  
}
```

```
// --size=<N> to specify different value on command line  
config const size : int = 10;
```

```
// two arrays of same size with different indexing
```

```
var B : [1..size, 1..size] real = 1.0;
```

```
var C : [0..<size, 0..<size] real = 1.0;
```

```
// call function directly on arrays
```

```
writeln(frobenius_norm(B));
```

```
writeln(frobenius_norm(C));
```

Chapel: serial language features

- iterators and zipping

○ ○ ○

```
use IO.FormattedIO;
```

```
var ordinals = ('first','second','third');
```

```
for (ord,i) in zip(ordinals,1..) {  
    writeln('ordinal "%s" corresponds to number "%i"'.format(ord,i));  
}
```


Chapel: serial language features

- Also:
 - multiple outputs
 - default arguments
 - objects

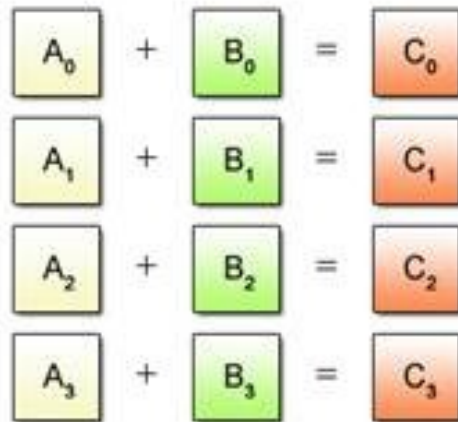
○ ○ ○

```
proc first_four_powers(a) {  
    var p1 = a;  
    var p2 = p1*a;  
    var p3 = p2*a;  
    var p4 = p3*a;  
    return (p1, p2, p3, p4);  
}  
  
var a : int = 2;  
writeln('The first four powers of ',a,' are:');  
for p in first_four_powers(a) {  
    writeln(' ',p);  
}
```

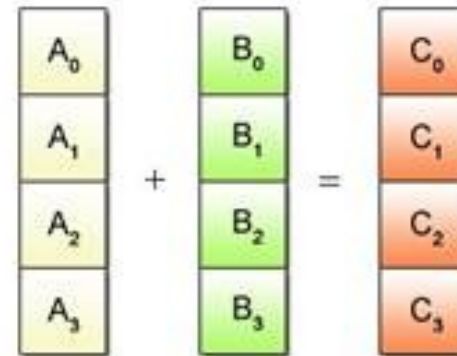
Parallel programming models: SIMD

- single instruction multiple data
 - gives up flexibility for speed
 - CPU instructions like addpd
 - contributes to performance of Numpy* and GPUs

(a) Scalar Operation



(b) SIMD Operation



source

Parallel programming models: SPMD

- single program multiple data

- trade flexibility for program complexity
- same program runs on every node, doing possibly different things simultaneously
- MPI is dominant standard in HPC, designed for distributed memory

```
# send vector in chunks to other nodes
if (node == 0):
    for i in [0,...,n_nodes]:
        MPI_SEND(a[8*i:8*(i+1)])
        MPI_SEND(b[8*i:8*(i+1)])
```

```
# all nodes receive chunks
MPI_RECEIVE(a_local)
MPI_RECEIVE(b_local)
c_local = a_local + b_local
```

```
if (node == 12):
    print("I'm node 12.")
```

```
...
```

Parallel programming models: directives

- Serial code annotated with directives for compiler
 - OpenMP and OpenACC both work for CPUs and GPUs
 - preserves portability (AMD vs Intel, Nvidia vs AMD)
 - often combined with MPI

```
#pragma omp parallel for
for(int i=0; i<ARRAY_SIZE; i++)
    c[i] = a[i] + b[i];
}
```

Note: “directives” summarizes a way of interacting with source code, but the directives themselves can do many things, including SIMD on CPUs or GPU kernels

Parallel programming models: PGAS

- Partitioned Global Address Space
 - global address space: threads can read/write remote data
 - partitioned: data designated as local versus global
 - examples: Unified Parallel C, CoArray Fortran, **Chapel**

MPI: isolated processes with isolated memories exchange messages

OpenMP: multiple threads can read and write a shared memory



PGAS

- virtually global address space realized across distributed hardware with non-uniform memory access times
- compiler uses program as specification to map processes to hardware

[source 1](#)

[source 2](#)

Chapel: locales

- Resources with compute and memory
 - most naturally a (many core) compute node
- Hierarchical:
 - may contain sublocales (like GPUs)

○ ○ ○

```
for loc in Locales {  
  writeln("Started on locale #", here.id);  
  on loc do  
    writeln("Hello from locale #", here.id);  
    writeln("Back on locale #", here.id, "\n");  
}
```

***this is not a parallel program**

Chapel: task parallelism

- begin:
 - fork one non-blocking process
- cobegin
 - fork one process per statement, blocking
- coforall
 - fork one process per iteration, blocking

○ ○ ○

```
writeln('Hello 1 (always executes first)');
cobegin {
    writeln('Hello 2');
    writeln('Hello 3');
    writeln('Hello 4');
}
writeln('Hello 5 (always executes last)');
```

○ ○ ○

```
writeln('Hello 1 (always executes first)');
coforall i in 2..4 {
    writeln('Hello ',i);
}
writeln('Hello 5 (always executes last)');
```

Chapel: data parallelism

- forall:
 - behavior dictated by supplied iterators, which determine degree of concurrency and location of execution
 - requires order-independence of loop body

○ ○ ○

```
// this is a data race  
forall i in 2..n-1 do  
  A[i] = A[i-1] + A[i+1];
```


Chapel: orthogonality of locality and parallelism

○○○

```
// serial
for loc in Locales do
  on loc do
    writeln("Hello from locale ", loc.id, " of ", numLocales);

// parallel
coforall loc in Locales do
  on loc do
    writeln("Hello from locale ", loc.id, " of ", numLocales);
```

Let's run some code!

- Code in chapel-introduction/src/parallel/triad
- Compile a Chapel source file with

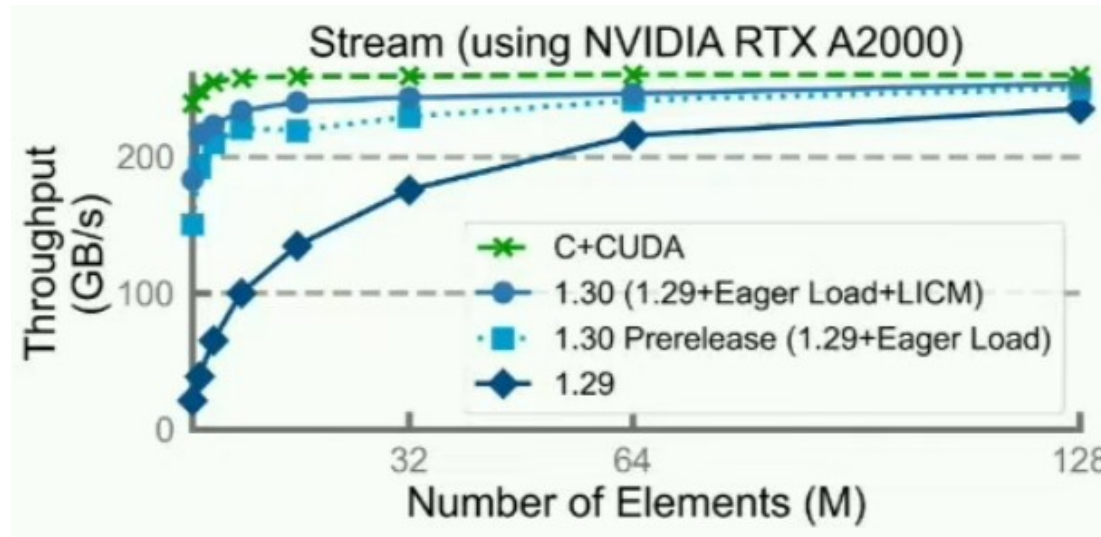
```
○ ○ ○
```

```
$ chpl --fast -o build/<name> <name>.chpl
```

- Compile all Fortran source files with make
- Setup times: 0.58 s (Fortran), 0.41 s (Chapel)

Chapel is ready for prime time, but...

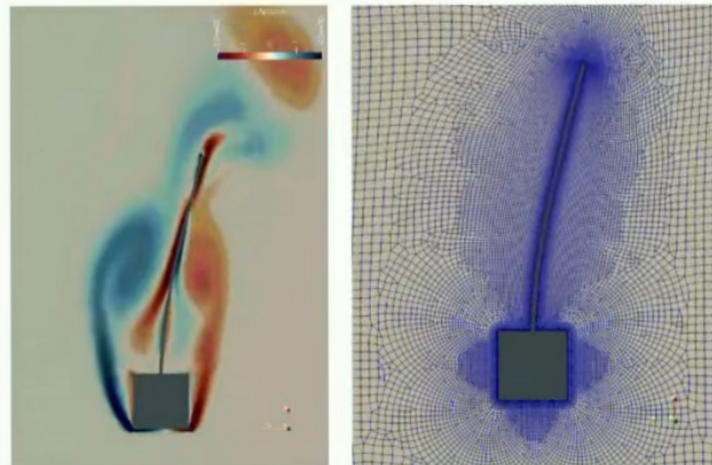
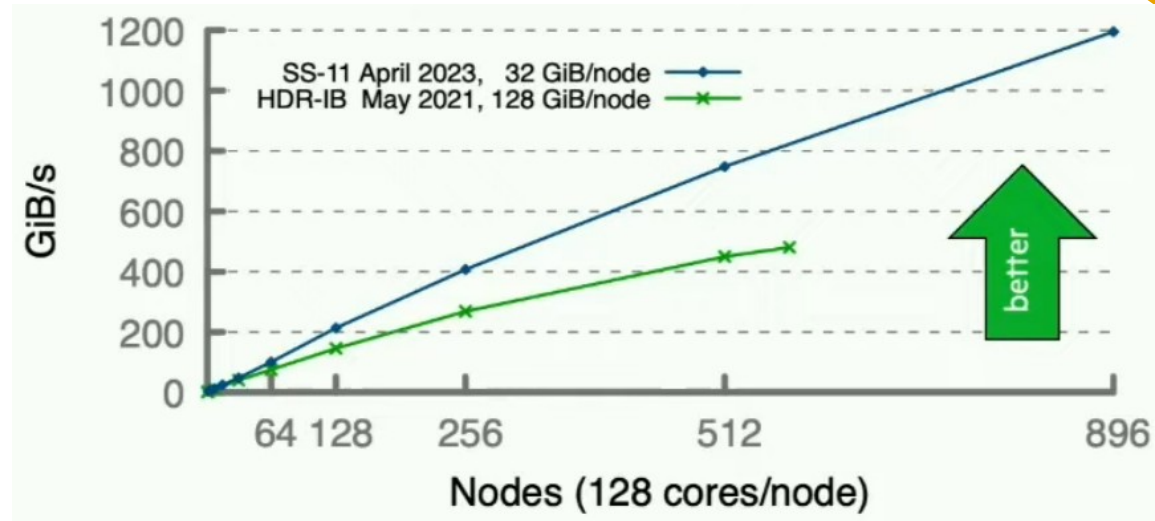
- compile times
- tooling
- hardware*
- packaging
- availability on HPC systems



source

Success stories

- Sorting:
 - (near) world record performance in 100 lines of Chapel
- CHAMPS
 - 3D unstructured CFD, airplane simulation
 - “students take 3 months to do what used to take 2 years”



source

Wrapping up

- Could not cover:
 - language interoperability: C, Fortran, Python
 - iterators, ranges, (sparse) domains, distributions, layouts, atomics, synchronization
 - GPUs
- Resources:
 - [Learn Chapel in Y minutes](#)
 - [Official Cray/HPE talk on Chapel](#)
 - [Online compiler](#)